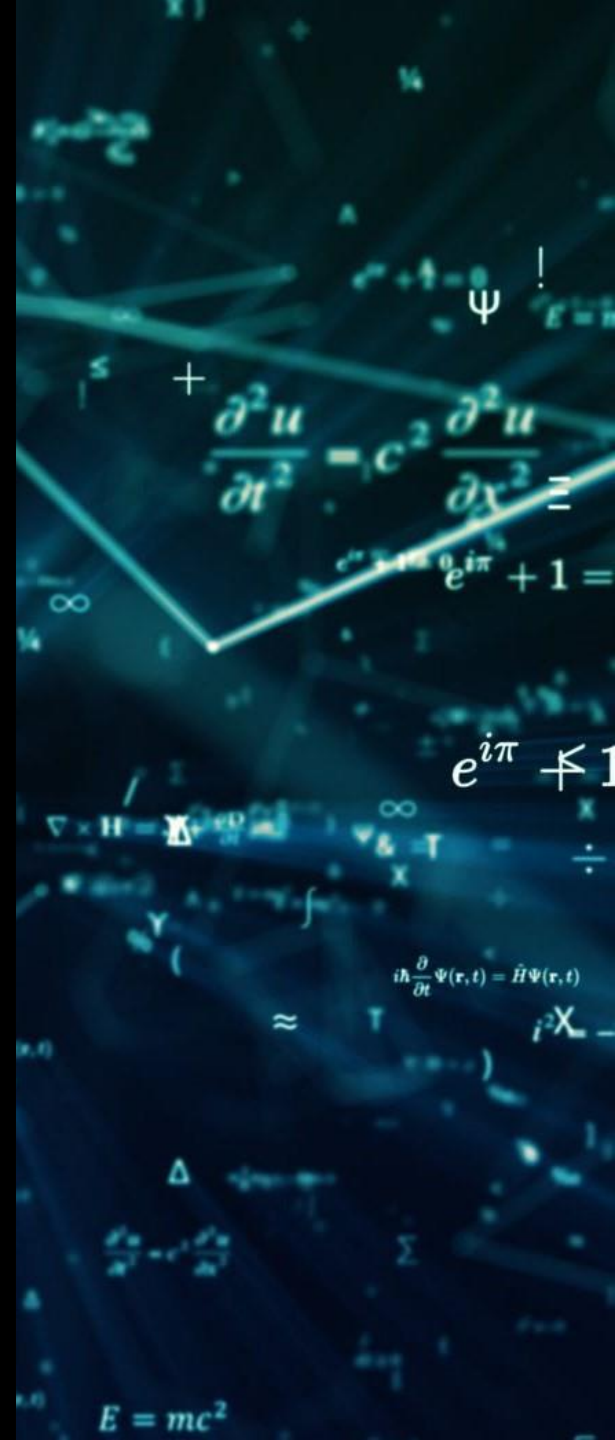


Artificial Intelligence Algorithms and Mathematics

CSCN 8000



Classification

- Solutions for Overfitting
- PCA
- FDA/LDA



Ways to reduce overfitting




- Overfitting occurs when a model learns not only the underlying patterns in the training data but also the noise and random fluctuations
- To reduce the overfitting, a general rule of thumb is to try to decrease the model complexity, but not too much to avoid underfitting.
- Specifically, techniques that can be used include:
 - **Cross-Validation:** Assess the model's performance on multiple subsets of the data to ensure good generalization to different portions of the dataset.
 - **Increase Dataset size:** A larger dataset can provide more diverse examples and help the model generalize better.
 - **Early Stopping:** Stop training once the performance on the validation set starts to degrade, preventing the model from overfitting the training data.
 - **Dropout:** random neurons are "dropped out" (ignored) during each training iteration. This helps prevent co-adaptation of neurons and reduces overfitting.
 - **Data Augmentation:** Increase the size of the training dataset by applying various transformations to the existing data
 - **Regularization?**

Regularization



- The intuition is that one way to decrease the overall model complexity is to reduce the magnitude of the learned weights \vec{w} . This way we won't be assigning very high weights to certain features over the others → Simpler model
- This could be achieved by adding a regularization term to the loss function of any machine learning algorithm, such that,
$$L_{reg}(\vec{w}, \mathbf{b}) = L(\vec{w}, \mathbf{b}) + [\beta * \text{Regularization Term}]$$
- This will allow the model to not only prioritize minimizing the cost term, but also minimizing the magnitude of the weights (controlled by the hyperparameter β)
- There are different types of regularization terms that could be used.




L1 Regularization (Lasso)



- L1 regularization introduces a penalty term proportional to the absolute values of the model parameters.
- Mathematically,

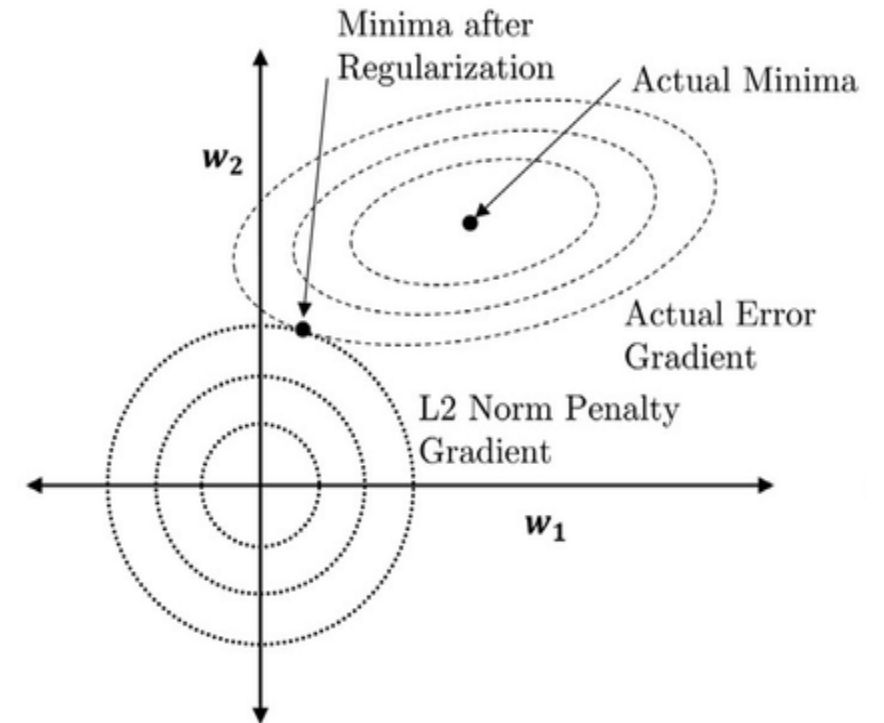
$$L_{reg}(\vec{w}, \mathbf{b}) = L(\vec{w}, \mathbf{b}) + \left[\beta * \sum_p |\vec{w}_p| \right]$$

- This encourages the model to prefer a sparse set of features, effectively driving some of them to exactly zero.
- Not only does it help prevent overfitting but also performs feature selection by excluding less relevant features.



L2 Regularization (Ridge)


- L2 regularization introduces a penalty term proportional to the square of the model parameters.
- Mathematically,
$$L_{reg}(\vec{w}, \mathbf{b}) = L(\vec{w}, \mathbf{b}) + \left[\beta * \sum_p (\vec{w}_p)^2 \right]$$
- This discourages the weights from becoming too large, preventing individual features from dominating the model.
- It distributes the importance of features more evenly, leading to a smoother model.



Dimensionality Reduction




- The efficiency of ML methods depends crucially on the choice of features that are used to characterize data points.
- Target → have a small number of highly relevant features to characterize data points.
- Dimensionality Reduction techniques reduce the number of input variables or features in a dataset while retaining its essential characteristics
- Benefits of dimensionality reduction:
 - Reduce excessive resource requirements
 - Reduce the probability of overfitting
 - Make data visualizations easier



Principal Component Analysis (PCA)



- Intuitively, if we want to project the following 2D points to only one dimension, which one of the following best preserves the shape of the dataset?




- Generally, it looks like in the best dimension, the projected data points have the largest possible variance.

Principal Component Analysis (PCA)



- The PCA is a dimensionality reduction technique that transforms high-dimensional data into a new coordinate system (principal components/axes).
- Target → **Maximize** the variance of projected data.
- Doesn't need any class labels (unsupervised)
- The projected data must obey certain properties:
 - Each new feature (principal component/axis) is a linear combination of the original features (axes)
 - Each principal component (axis) must be perpendicular to all other principal components (axes).
 - Each principal component (axis) must be a unit vector (magnitude = 1)



Principal Component Analysis (PCA)



- **Covariance Matrix:** provides important information about the relationships between pairs of variables in a dataset.
- Given a dataset $X \in R^{D*N}$ with D features and N rows, the covariance matrix will have size $D * D$ can be computed as follows:

$$C(X) = \frac{1}{N} (X - \bar{X})(X - \bar{X})^T$$

- \bar{X} represents the mean of each feature in X .
- Cell C_{ij} represents the co-variance between feature i and feature j .
- If one feature increases/decreases and the other feature increases/decreases at the same time $\rightarrow C_{ij} \neq 0$.
- If the features are not correlated $\rightarrow C_{ij} = 0$.

Eigen Vectors and Eigen Values



- Given a vector v , apply linear transformation with square matrix A .

$$A \cdot v = ?$$

$$A \cdot v = \lambda v \rightarrow [\text{Special Case}]$$

$$A = \begin{bmatrix} 1 & 2 \\ 2 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 3 \\ 3 \end{bmatrix} = 3 \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} -1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ -1 \end{bmatrix} = -1 \begin{bmatrix} -1 \\ 1 \end{bmatrix}$$

$$\lambda_1 = 3, \quad v_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$$

$$\lambda_2 = -1, \quad v_2 = \begin{bmatrix} -1 \\ 1 \end{bmatrix}$$

Principal Component Analysis (PCA)



- Given the initial dataset with D features, we want to get the first axis of the new coordinate space.
- We will apply a linear transformation $u \in R^{D*1}$ on each original data point $x \in R^{D*1}$, such that the value z of the projected point at the new axis is formulated as,

$$z = u^T x$$

- Applying the same equation on the full dataset $X \in R^{D*N}$, such that,

$$Z = u^T X$$

- Target → **Maximize** the variance of projected data (1D) → **Maximize** the covariance matrix of projected data (Multi-Dimensional)

Principal Component Analysis (PCA)



- Target → **Maximize** the covariance matrix of projected data
- $Cov(Z) = \frac{1}{N} (Z - \bar{Z})(Z - \bar{Z})^T$
- $Cov(Z) = \frac{1}{N} (u^T X - u^T \bar{X})(u^T X - u^T \bar{X})^T$
- $Cov(Z) = \frac{1}{N} \left(u^T (X - \bar{X}) \right) \left(u^T (X - \bar{X}) \right)^T \rightarrow \text{Expand 2}^{\text{nd}} \text{Transpose}$
- $Cov(Z) = \frac{1}{N} u^T (X - \bar{X})(X - \bar{X})^T u = u^T \left[\frac{1}{N} (X - \bar{X})(X - \bar{X})^T \right] u$
- $Cov(Z) = u^T S u$, where $S = Cov(X)$

Principal Component Analysis (PCA)



$$\begin{aligned} & \text{maximize } u^T S u, \\ & \text{such that } u^T u = 1 \end{aligned}$$

- The constraint guarantees that the axis u is a unit vector.
- By borrowing the concept of Lagrangian Multipliers, we will translate the equation to a loss function to be **minimized**, such that,

$$L(u, \lambda) = -(u^T S u - \lambda(u^T u - 1))$$

- To minimize the loss with respect to $u \rightarrow \text{Solve } \frac{dL}{du} = 0$

Matrix/Vector Rules and Derivatives



Rule	Comments
$(\mathbf{AB})^T = \mathbf{B}^T \mathbf{A}^T$ $(\mathbf{a}^T \mathbf{B} \mathbf{c})^T = \mathbf{c}^T \mathbf{B}^T \mathbf{a}$ $\mathbf{a}^T \mathbf{b} = \mathbf{b}^T \mathbf{a}$	order is reversed, everything is transposed as above (the result is a scalar, and the transpose of a scalar is itself)
$(\mathbf{A} + \mathbf{B})\mathbf{C} = \mathbf{AC} + \mathbf{BC}$ $(\mathbf{a} + \mathbf{b})^T \mathbf{C} = \mathbf{a}^T \mathbf{C} + \mathbf{b}^T \mathbf{C}$	multiplication is distributive as above, with vectors
$\mathbf{AB} \neq \mathbf{BA}$	multiplication is not commutative

Scalar derivative	Vector derivative
$f(x) \rightarrow \frac{df}{dx}$	$f(\mathbf{x}) \rightarrow \frac{df}{d\mathbf{x}}$
$bx \rightarrow b$	$\mathbf{x}^T \mathbf{B} \rightarrow \mathbf{B}$
$bx \rightarrow b$	$\mathbf{x}^T \mathbf{b} \rightarrow \mathbf{b}$
$x^2 \rightarrow 2x$	$\mathbf{x}^T \mathbf{x} \rightarrow 2\mathbf{x}$
$bx^2 \rightarrow 2bx$	$\mathbf{x}^T \mathbf{B} \mathbf{x} \rightarrow 2\mathbf{B}\mathbf{x}$

Principal Component Analysis (PCA)



$$L(u, \lambda) = -(u^T S u - \lambda(u^T u - 1))$$

- To minimize the loss with respect to $u \rightarrow$ Solve $\frac{dL}{du} = 0$
- $\frac{dL}{du} = -(2Su - 2\lambda u) = 0$

$$Su = \lambda u \rightarrow \text{Result of } \frac{dl}{du} = 0$$

- We reach a formulation exactly similar to the one of eigenvalues and eigenvectors.
- In other words, u is considered an eigenvector of the covariance matrix S of the original data and λ is the associated eigenvalue.



Principal Component Analysis (PCA)



$$Su = \lambda u$$

- To project original data $X \in R^{D \times N}$ to P features, where $P \leq D$:
 - Calculate the Covariance Matrix $S = \frac{1}{N} (X - \bar{X})(X - \bar{X})^T$
 - Get all the possible eigenvalues and eigenvectors of S .
 - Sort the eigenvalues in descending order:
 - The Largest eigenvalue corresponds to the (eigenvector) axis with highest variance of data projected on that axis.
 - Lower eigenvalues correspond to axes that are worse in preserving the characteristics of the data.
 - Get the highest P eigenvalues and their eigenvectors.
 - Construct full matrix $U \in R^{P \times D}$ by stacking all chosen eigenvectors vertically (row-wise)
 - Get the final full projected dataset $Z \in R^{P \times N} \rightarrow Z = UX^T$


Principal Component Analysis (PCA)



Fisher Discriminant Analysis (FDA)




- Assume we want to project our features to fewer dimensions while maintaining the separability of our classes.
- A possible approach would be to **maximize** the distance between the centers (means) of the projected classes.
- In the following example, does the proposed axis maintain the best separability between the classes?





Fisher Discriminant Analysis (FDA)



- It looks like maximizing the distance between means of projected classes is not enough if the classes are wide-spread with high variance.
- An additional constraint could be imposed by minimizing the within-class variance of the projected data.
- Combining the two constraints leads to a new axis (dimensionality) that maintains the separability of the original data with minimum or no overlap.



Fisher Discriminant Analysis (FDA)



Fisher Discriminant Analysis (FDA)



- Fisher's Linear Discriminant Analysis (FDA) is a linear dimensionality reduction technique that aims to project high-dimensional data into a lower-dimensional space while maximizing the separation between classes.
- This is achieved through two targets:
 - **T1: Maximizing Inter-Class Variance (Distance between class means)**
 - **T2: Minimizing the Intra-Class Variance (Within-class variance).**

Fisher Discriminant Analysis (FDA)



- Assume that we have two classes to be projected.
- We will apply a linear transformation $u \in R^{D*1}$ on each original data point $x_{\{0,1\}} \in R^{D*1}$ belonging to classes 0 *and* 1, such that the value z of the projected point at the new axis is formulated as,

$$z = u^T X$$

- The means of the points in each class are formulated as,

- $\mu_0 = \frac{1}{N_0} \sum_{i=0}^{N_0} x_0^i,$

- $\mu_1 = \frac{1}{N_1} \sum_{i=0}^{N_1} x_1^i$

Fisher Discriminant Analysis (FDA)



- $\mu_0 = \frac{1}{N_0} \sum_{i=0}^{N_0} x_0^i, \quad \mu_1 = \frac{1}{N_1} \sum_{i=0}^{N_1} x_1^i$
- **For Target 1:** Distance between the projected means is formulated as:
 - $(u^T \mu_0 - u^T \mu_1)^2 = (u^T \mu_0 - u^T \mu_1)^T (u^T \mu_0 - u^T \mu_1)$
 - $(u^T \mu_0 - u^T \mu_1)^2 = (\mu_0 - \mu_1)^T u u^T (\mu_0 - \mu_1)$
 - $(u^T \mu_0 - u^T \mu_1)^2 = u^T (\mu_0 - \mu_1) (\mu_0 - \mu_1)^T u$
 - $(u^T \mu_0 - u^T \mu_1)^2 = \mathbf{u}^T \mathbf{S}_B \mathbf{u}, \quad S_B = (\mu_0 - \mu_1)(\mu_0 - \mu_1)^T$
 - Where S_B represents the distance between class means before projection.

Fisher Discriminant Analysis (FDA)



- Recall from PCA $\rightarrow Cov(Z) = u^T S u$, where $S = Cov(X)$
- **For Target 2:** Within class-variance for the two classes can be formulated as:
 - $Cov(Z_0 + Z_1) = Cov(Z_0) + Cov(Z_1)$
 - $Cov(Z_0 + Z_1) = u^T S_0 u + u^T S_1 u$, where $S_0 = Cov(X_0), S_1 = Cov(X_1)$
 - $Cov(Z_0 + Z_1) = u^T (S_0 + S_1) u = \mathbf{u}^T \mathbf{S}_W \mathbf{u}$, $S_W = S_0 + S_1$
 - Where S_W represents the within-class variance of the two classes altogether.

Fisher Discriminant Analysis (FDA)



- To Achieve both **Target 1** and **Target 2**, our target could be formulated as follows:

$$\text{maximize } \frac{u^T S_B u}{u^T S_W u}$$

- Recall that our new axis needs to be unit vector (from PCA), to enforce it we can formulate the target as follows:

$$\text{maximize } u^T S_B u, \quad \text{s.t. } u^T S_W u = 1$$

- To formulate it as a loss function, we will borrow the concepts from Lagrangian Multipliers:

$$L(u, \lambda) = - \left(u^T S_B u - \lambda (u^T S_W u - 1) \right)$$

- The negative sign is added to minimize rather than maximize

Fisher Discriminant Analysis (FDA)



$$L(u, \lambda) = - \left(u^T S_B u - \lambda (u^T S_W u - 1) \right)$$



- To minimize the loss with respect to $u \rightarrow$ Solve $\frac{dL}{du} = 0$
- $\frac{dL}{du} = 2S_B u - 2\lambda S_W u = 0$
- $S_B u = \lambda S_W u \rightarrow [S_W^{-1} S_B] u = \lambda u$
- We reach a formulation exactly similar to the one of eigenvalues and eigenvectors.
- In other words, u is considered an eigenvector of the matrix $S_W^{-1} S_B$ calculated from the original data and λ is the associated eigenvalue.
- To transform the full dataset, follow the same steps as PCA, but with calculating $S_W^{-1} S_B$ in the first step instead.

LDA vs FDA



- Both Linear Discriminant Analysis (LDA) and FDA refer to the same technique which aims to project the data to lower dimensions while maximizing the class separability.
- LDA is the direct extension of FDA to work with two **or more** classes.
- LDA is not only doing dimensionality reduction, but also computes the linear decision boundary between the classes in the projected space.
- LDA makes important assumptions about the shape of the data:
 - All classes follow a gaussian (normal) distribution
 - All classes have equal (identical) covariance matrices.
- If any of those assumptions doesn't hold, LDA won't perform well in classification or dimensionality reduction.

Linear Discriminant Analysis



Extra Resources



- More information regarding the Matrix/Vector Derivatives:
 - <https://www.math.uwaterloo.ca/~hwolkowi/matrixcookbook.pdf>

Thank you!



- Any questions?



Disclaimer



Due to nature of the course, various materials have compiled from different open source resources with some moderation. I sincerely acknowledge their hard work and contribution



Thank You

Youssef Abdelkareem

yabdelkareem@conestogac.on.ca